

JEET DAVE

Jersey City, NJ | 302-312-3276 | jeetd6220@gmail.com | [LinkedIn](#) | [Portfolio](#)

PROFILE

Data engineering-focused MS Data Science student with hands-on experience designing ETL data pipelines, data modeling, and large-scale analytics — analyzing data and creating reports using SQL and visualization tools (Tableau, Power BI). Skilled in Python, big data tools, and distributed computing (Apache Spark, Airflow, Kubernetes), supporting ML workflows through feature engineering on non-relational databases, with structured thinking to break down ambiguous problems and a passion for turning large, complex data sets into insights that drive business decisions.

PROFESSIONAL EXPERIENCE

NEXGITS PRIVATE LIMITED

Ahmedabad, India

Data Analyst Intern

Jun 2024 to July 2025

- Dove into data and processes to identify operational challenges and recommend actionable solutions; designed ETL pipelines in Python and SQL to extract, cleanse, and transform structured data across multiple relational tables, identifying and resolving data inconsistencies, reducing prep time by 30% and surfacing insights that drove real business change.
- Defined and tracked KPIs critical to business unit goals, translating data insights into clear, actionable recommendations for senior stakeholders; built interactive Power BI, Tableau, and Excel dashboards and standardized reports communicating key trends that directly informed 3 strategic decisions within the first month.
- Partnered across teams to operate independently in a fast-paced, ambiguous environment and deliver results on tight timelines; automated analytical workflows using Python on AWS, reducing manual effort by 40% and improving the speed and reliability of reporting.

ACADEMIC & PERSONAL PROJECTS

Kubernetes Data Platform — Distributed Data Processing on AWS EKS

May 2026 to Present

- Designing a production-grade data platform on Kubernetes (AWS EKS) processing BTS flight-delay data through a Bronze/Silver/Gold **medallion architecture**, with infrastructure fully provisioned via Terraform (EKS cluster, managed node groups).
- Deploying and orchestrating **Apache Spark and Apache Airflow** on Kubernetes via Helm, using AI-assisted debugging and documentation to work through cluster autoscaling, IRSA permissioning, and Helm chart migration issues.

FraudSentinel — Real-Time Fraud Detection Platform

Mar 2026 to Apr 2026

- Architected an end-to-end fraud detection platform across 10 AWS services (SQS, Lambda, DynamoDB, S3, Glue, Athena, Step Functions, SNS, SageMaker, CloudWatch) processing 507K+ PaySim transactions with a 7.7% flagging rate and zero dead-letter-queue failures across all runs.
- Built a real-time Lambda fraud processor (P99 latency: 761ms) applying 4 detection rules against a **DynamoDB TTL-based feature store** — performing **feature engineering** on velocity-tracking signals from raw transaction data — with SNS alerts delivering structured fraud notifications in under 30 seconds to support fraud-loss prevention at scale.
- Designed a Bronze-Silver-Gold medallion data lake using AWS Glue PySpark ETL jobs with automated catalog registration and Athena SQL analytics, and trained an XGBoost fraud classifier (AUC: 1.00, F1: 1.00) deployed to a SageMaker Serverless Inference endpoint — orchestrated via Step Functions and EventBridge for \$0.13 total compute cost.

Retail Analytics Pipeline — ETL & Data Warehouse Design

Oct 2025 to Nov 2025

- Designed and built a cloud data warehouse using **Snowflake, Snowpark Python, and dbt Core** — **6 data models** with **52 passing data-quality tests** — transforming raw structured retail data into analytics-ready tables.
- Built ETL logic to handle a mix of structured sources end-to-end, with a Streamlit dashboard layered on top for business-facing reporting.
- Optimized SQL transformations and dbt models to improve data quality, maintainability, and downstream analytics performance — enabling faster, more reliable reporting for business decision-making.

E-Commerce Data Pipeline — EDA, Data Modeling & ML

Feb 2026 to Apr 2026

- Built an end-to-end pipeline handling **100K+ records** using Python, SQL, and MySQL, designing relational **data models** and performing exploratory data analysis to surface KPI-level insights across the business.
- Built and evaluated 3 ML models — Sales Forecasting (Linear Regression, R^2 : 0.65), Customer Churn Prediction (63% accuracy), and K-Means Segmentation — redefining customer targeting across 4 segments (VIP, Regular, At-Risk, Dormant).
- Delivered a scalable interactive dashboard in **Tableau** and Plotly with live KPI filters, translating complex data into clear, executive-ready insights for non-analytical stakeholders.

TECHNICAL SKILLS

- **Programming Languages:** Python, SQL, C++, C
- **Data Pipelines & Data Modeling:** ETL/ELT Pipelines, Data Modeling, Dimensional Modeling, Schema Design, Medallion Architecture, Data Quality Testing
- **Big Data & Distributed Computing:** Apache Spark, PySpark, Apache Airflow, Kubernetes, Docker, Hadoop (MapReduce, Distributed Storage)
- **Databases:** MySQL, PostgreSQL, Snowflake, Non-Relational Databases (DynamoDB / NoSQL)
- **Analysis, Reporting & Visualization:** SQL Querying, Tableau, Power BI, Excel, Streamlit, Feature Engineering, Model Automation
- **Cloud & Infrastructure:** AWS (S3, Lambda, EKS, Glue, Step Functions, SageMaker, CloudWatch, DynamoDB), Terraform, Git/GitHub
- **Collaboration & Problem-Solving:** Structured Thinking, Stakeholder Management, Cross-Functional Collaboration, AI-Assisted Development

EDUCATION

New Jersey Institute of Technology

Newark, NJ

MS in Data Science | GPA: 3.5

Expected May 2027

Indus University

Ahmedabad, India

BS in Computer Science & Engineering | GPA: 3.8

May 2025

CERTIFICATIONS

- AWS Machine Learning | Amazon Web Services | Jul 2024 to Nov 2024